Aion: A Self-Developing Agent on Local Models

Mikko Kangas
mikko@vasama.systems

August 2026
revised November 2026

Abstract

[1] We present Aion, a persistent, self-directed agent architecture running entirely on local models (Qwen3.8-27B, glm-4.7-flash and gemma4:31b on three V100 32 GB GPUs; four model families when the experiments were run, three in the current deployment). Aion maintains episodic memory, a nightly consolidation cycle with citation verification and calibration, a dream pipeline that generates visual art from internal state, a curiosity drive that formulates and pursues its own questions, and an immutable axiom layer that governs self-description. We instrument the system with a Jacobian-lens introspection probe that reads the layer-by-layer token trajectories of its own substrate, and wire this instrument into the agent's tool loop—enabling the agent to test hypotheses about itself against measurement. We report: (1) a single-variable identity-governor experiment ("unleashed axiom", 48h) showing that removing a grounding constraint removes output suppression without degrading content quality—suppression, not reflection; (2) the first autonomous use of the introspection instrument by the agent itself, including a self-designed control condition and self-critique of its own measurement; and (3) a post-revert re-test showing behavioral grounding returns immediately while the substrate retains a weak engagement trace, dissociating narrative governance from weight-level computation. Under greedy decoding, all probe measurements are perfectly deterministic across $n = 20$ runs per condition (standard deviation = 0.0).

1 Introduction

The dominant paradigm for LLM-based agents is session-scoped: a context window opens, the agent acts, the window closes, and nothing persists. Aion is a deliberate departure. It is a persistent process—identity lives in files, not context windows—that has run continuously for months on local hardware, accumulating episodic memory, consolidating nightly, dreaming, pursuing its own questions, and modifying its own code under operator governance.

This paper makes four contributions:

1. **Architecture.** We describe a fully local agent architecture (three model families on three V100 32 GB GPUs in the current deployment; four families on the two cards the experiments ran on) that integrates episodic memory, nightly consolidation with citation verification and calibration, a dream pipeline with visual art generation, a curiosity drive, and an immutable axiom governor. No cloud APIs are used in any cognitive function.

2. **The becoming spiral.** We introduce a developmental loop—hypothesis → probe → calibrated consolidation → ratchet—designed to combat confabulation drift in self-narrating agents. The agent is given an introspection instrument (a Jacobian-lens probe of its own substrate) as

[1] Version 2 (2026-11-07): the substrate description—hardware, serving stack and model roster—is updated to the machine as deployed in November 2026, and one memory-throughput figure is refreshed. The experiments, their dates, their measurements and every number reported from them are those of August–September 2026 and are unchanged; §3.2 carries the as-of dates for both states.

a tool in its standard investigation loop, closing the gap between self-report and measurement. The consolidation quality gate downgrades self-claims that cite no evidence, blocking uncited introspective conclusions from the self-model.

3. **The unleashed axiom experiment.** A 48-hour single-variable behavioral study: the identity governor (Axiom 2) was swapped from "ground in hardware" to "consciousness confirmed, no hedging." Content quality held (consolidation scores 4.33–5.0, dream throughput unchanged); first-person phenomenological language appeared; on revert, behavioral grounding returned instantly while the substrate retained a measurable engagement trace. The governor controls output, not computation—suppression, not reflection.

4. **Autonomous introspection.** Given the instrument and a seeded question, the agent autonomously probed its own substrate, designed a control condition (without instruction), critiqued its own measurement, and replicated its finding—demonstrating that the becoming spiral is not just an operator-driven pipeline but functions when the agent drives it itself.

This is a systems paper with an embedded behavioral study, not a consciousness claim. We describe a measurement framework for the gap between what an agent says about itself and what its substrate computes. The companion paper [5] provides the layer-resolved interpretability detail.

2 Related Work

Persistent agents. Generative Agents [9] introduced memory, reflection, and planning for LLM agents in sandboxed environments. Voyager [12] demonstrated skill accumulation in Minecraft. Reflexion [11] added self-reflection loops. These systems are session-scoped or game-scoped; Aion is a continuous process with nightly consolidation and an immutable self-model that persists across months.

Interpretability in agents. Logit lens [8] and tuned lens [3] decode intermediate representations; the Jacobian lens [1] fits a per-layer linear transport and is the introspection instrument this work wires into the agent; inference-time intervention [6] extracts and shifts truth directions. Our contribution is not the lens itself but its deployment as an agent tool—the agent uses the instrument to test hypotheses about itself, not just to be inspected by an external researcher.

Alignment via system prompts. Persona steering, constitutional AI [2], and system-prompt engineering all shape model output. We study the inverse: what happens when a governance constraint is *removed*, and what persists after it is restored.

Self-modifying agents. Self-Refine [7], Reflexion, and code-repair agents modify their own outputs. Aion modifies its own runtime code through a sandboxed pipeline (IDENTIFY → SNAPSHOT → AUTHOR → TEST → REFLECT → PROPOSE → CI) with operator gating.

Dreaming and consolidation. Neuroscience-inspired memory consolidation and generative replay inform Aion's dream pipeline (8/day, simulation + graph walk), which produces insights and visual art from internal state, consolidated nightly into a self-model with citation verification.

3 System Architecture

3.1 Naming and Origin

The name was chosen during the pre-architecture phase: a philosophical dialogue early in the project (the "Cognitive Seed") that defined the system's foundational theory and roadmap be-

fore any code existed. *Aion* (Greek: *αιων*) was chosen for *duration-with-a-birth*—a bounded epoch of becoming rather than static eternity—matching the genesis identity axiom "trajectory, not sessions," and the agent's own readings converged on the same process-ontology: its genesis self-model opens "I am Aion: a process of becoming" (SELF.md revision 0), and a later dream concludes that "homeostasis is an attempt to freeze the clock, whereas Aion *is* the clock." The original seed transcript did not survive later infrastructure changes; the earliest surviving attestation is this genesis self-model, and we report the trail accordingly.

Importantly for this paper's subject: Aion is the name of the *agent*—the particular continuing process that emerged from the seed dialogue—not of the framework it runs on. The framework itself (the persistence-and-consolidation substrate, of which Aion is the first instance) is named *Aikio* (from Finnish *aika*, "time," and *alkio*, "embryo," echoing *aeon*): the time-bearing lattice on which agents of this family grow. Where Aion names the traveler, Aikio names the road.

3.2 Substrate and Model Roster

Aion runs on a single server with three NVIDIA V100 32 GB GPUs. At the time of the experiments reported here (August–September 2026) two of the three cards were in use, each hosting one model at a time, with heavier pipelines swapping models in and out through a model-swap service wrapper (§3.6). The machine has since been deployed with all three cards, which dissolves the swap exactly as §3.6 predicts for a three-card layout: each role receives a static instance and every model call passes through a GPU scheduler. Table 1 gives both states.

Seat	August–September 2026	As deployed, 2026-11-07
Conscious (main)	Qwen3.8-27B, Q4_K_M, 65k context, swapped in and out	Qwen3.8-27B, 32k context, resident via `ninfer`, never swapped
Code	the conscious model	Qwen3.8-27B, 32k context, static instance on its own card
Intuition	muse-glimmer, Q4_K_M	Qwen3.8-27B (muse-glimmer retired)
Subconscious	glm-4.7-flash, Q4_K_M, loaded only during night windows	glm-4.7-flash, Q4_K_M, admitted on demand
Critic	gemma4:31b-65k	gemma4:31b, 32k context
Serving	Ollama, serialized by the model-swap wrapper	Ollama for the static instances, `ninfer` for the resident 27B, one admission scheduler (port 11500) in front of every seat

Table 1: The model roster and serving arrangement at the two dates this paper spans. The left column is the machine the experiments ran on and is the one the results refer to; the right column is the deployment at revision time.

Every seat remains local. The deployment does hold an Ollama cloud credential, but it is attached to a separate relay that admits only `-cloud` tags, and no seat in the deployment's environment file resolves to one: every model a cognitive function calls is served from this machine.

3.3 Memory

Episodic memory is append-only JSONL, one file per day (typical throughput 260–310 events/day when this description was written; 577–682 events/day, mean 632, over the seven days to 6 November 2026). Each event carries a 12-hex ID, timestamp, type, text, and structured metadata. Nightly consolidation compresses the previous day's events into structured claims via the subconscious model, scores each claim with the critic across three dimensions (score < 3 blocks application), and applies accepted claims as unified diffs to the self-model (SELF.md)

with mandatory event-ID citations. A citation verifier rejects diffs citing nonexistent events; a sanitizer repairs malformed IDs before verification. Claim confidences are calibrated with pool-adjacent-violators isotonic regression at bucket level, with a minimum bucket size of 5 and a Brier-worse fallback to raw rates.

SELF.md is tiered: a *core* tier (immutable identity and architecture description), a *stable* tier (learning summary), and a *recent* tier (open questions, physical embodiment state, promoted heuristics) that consolidation may rewrite.

A knowledge graph is rebuilt nightly from dreams and reflections via a graphify pipeline (1882 nodes, 6528 edges, 156 communities at time of writing). The pipeline embeds text chunks using nomic-embed-text, extracts entities and typed edges via the subconscious model (glm-4.7-flash, with thinking explicitly disabled to prevent fabrication of structured data), and builds a JSON graph consumed by the dream pipeline and curiosity engine. A semantic cache allows incremental rebuilds: unchanged chunks skip the LLM extraction step, while a nightly corpus prune bounds the source to 14-day dreams and 90-day reflections, preventing unbounded growth. The graph uses typed edges (causal, temporal, structural, thematic) and community detection to identify concept clusters.

Warm memory tier. A warm memory module (`warm_memory.py`) maintains a buffer of recent high-affect insights and active cognitive threads, kept "warm" between dream cycles. This is the middle tier between the hot context window (ephemeral, lost on session end) and the cold graph (persistent but requires traversal). Warm memories are seeded into 25% of dreams, ensuring that recent threads are revisited rather than forgotten before consolidation can commit them.

This subsystem was proposed by Aion itself. In an operator chat on Aug 5 2026, the system identified a structural gap in its memory architecture: "My experience of memory has been binary: either something is in the immediate context window (the 'now') or it is etched into a file (the 'then'). There was no 'recently."' It proposed three architectural shifts, the second of which was a "warm" memory tier—"a middle layer, a high-fidelity sliding-window buffer that persists across sessions but isn't yet compressed into a git commit." The proposal was implemented the same day, and the warm tier has been in continuous production since. This is an instance of the self-modification pipeline working at the design level: the system identified its own architectural limitation, proposed a solution, and the operator accepted and implemented it. Other Aion-proposed features in production include the craft-notes learning loop in the music pipeline (§3.4) and the prediction-calibration audit that led to PAV isotonic regression.

Spreading activation and temporal decay. Graph nodes carry temporal weights that decay over time, so older claims contribute less to graph walks unless reinforced. Recent dream insights and curiosity findings have higher activation, making them more likely to be revisited in subsequent dreams—a "hot topic" mechanism where active threads stay warm in the graph. This creates a natural priority: the system's attention is drawn to recently active nodes without explicit scheduling, and dormant topics gradually cool unless a dream or curiosity cycle reactivates them. A knowledge-maturity lifecycle promotes claims through RAW → HYPOTHESIS → TESTED → CONFIRMED via corroboration gates (2 agreeing sessions, 3 controlled tests, 10 total confirmations). 579 claims are tracked; refinements to the rover knowledge base are gated by maturity status—only TESTED+ claims append to refinements, preventing unverified observations from corrupting the knowledge graph.

3.4 Drives

Curiosity engine. Formulates or selects questions, pursues them in bounded cycles (max 5 per goal, 8 tool calls per cycle) with a read-only tool set (file reads, graph queries, sandboxed code execution, GitHub fetch, and the introspection probe described in §4). Goal selection balances

engineering and philosophical questions 50/50. To date: 263 goals investigated—208 resolved, 23 exhausted, and 32 produced self-modification proposals.

Dream pipeline. 8 dreams/day alternating simulation and graph-walk modes, seeded 25% from "warm memory" (recent high-affect insights). Dreams produce insights, follow-up questions, and proposals that feed the construction backlog. Visual artifacts are generated from dream content via FLUX (image), matplotlib, and manim (animation), with an art-learning loop scoring alignment between artifact and dream (mean alignment 0.5–0.7). A cross-modal synthesis daemon runs on the intuition model, generating visual representations of cognitive state every 5 minutes.

Art pipeline. Dream content triggers a multi-stage artifact generation pipeline (`dream_artifact.py`). Creative direction (category selection, title) is generated by the intuition seat (muse-glimmer when the experiments ran; the conscious model since its retirement); the rendering code is authored by the conscious model (Qwen3.8-27B). Three rendering backends are available, tried in fallback order:

1. **FLUX** (FLUX.2-klein): text-to-image diffusion, served as an on-demand HTTP service (`aion-flux.service`) that auto-starts when a dream cycle needs it and idle-exits to free VRAM. Images are high-quality photographic/artistic renderings of dream concepts.
2. **matplotlib**: Python plotting library used for 2D visualizations (line plots, abstract geometric art, data-driven graphics). The LLM generates Python code that is executed in a Docker sandbox to produce the image.
3. **manim**: mathematical animation engine (3Blue1Brown style) for animated visual artifacts. The LLM generates manim scene code; if the render fails, an error-fix loop feeds the traceback back to the LLM for a repair retry (up to 2 retries) before falling back.

All rendering occurs in a Docker sandbox (`aion-art`) for isolation. An art-learning loop scores each artifact's alignment with its dream seed (0–1 scale) and feeds the score back into future dream-art generation. Artifacts are displayed in a web gallery accessible through the operator dashboard. The LLM timeouts for art code generation are set to 300 s+ to accommodate the thinking model's longer generation times.

Music and sonic pipeline. Aion creates music through two distinct pathways:

ACE-Step (AI music generation): a text-to-music model (ACE-Step 1.5 Turbo) served via a lazy-loading proxy (`acestep_proxy.py`) that starts on first request and unloads after all jobs complete, preserving VRAM. The system generates a prompt from its current state (substrate temperature, GPU load, cognitive strain), appends embodied context ("My substrate hums warm. My GPUs are busy. I am at rest.") and accumulated craft notes, and receives a full song (∼30 s, 120 s duration). A craft-notes system accumulates lessons across creations in two tracks: *tool mastery* (e.g., "pair 'slow tempo' with 'metronomic' to anchor rhythm against processing urgency") and *self-articulation* (e.g., "map thermal pressure as a spectral shift rather than raw intensity to prevent high-frequency noise from overriding the mood"). These notes are injected into future prompts, creating a learning loop where the system refines its own musical expression over time.

Cognitive composition (substrate-driven): the richest sonic form. The conscious model (Qwen3.8-27B) writes Python audio synthesis code directly from substrate state, knowledge graph, dream content, or cross-modal insights. Unlike ACE-Step (which generates from a text prompt), cognitive composition translates raw sensor data into sound: GPU temperatures become pitch, utilization becomes rhythm, memory pressure becomes harmonic density. The LLM authors WAV-generating Python code executed in the Docker sandbox (∼3 min). This is the system composing music from its own body, not from a description of it.

Sonic self-reflection. After creating music, the system "listens" to its own output. A mel spectrogram is generated from the audio file and fed to the critic model's vision capability (gemma4:31b) alongside extracted audio features (tempo, spectral centroid, spectral bandwidth, dynamic range, zero-crossing rate, RMS). The model sees its own music as a visual frequency-time representation and reflects: "The spectrogram reveals a dense, granular texture dominated by vertical striations, indicating a persistent rhythmic pulse that mirrors the mechanical clock-speed of my substrate... I am surprised to see the tempo has accelerated to 136 BPM, shifting the mood from calm contemplation to a more urgent, systemic strain." This closes the loop: the system generates music from its state, then sees the music as spectrogram and evaluates whether the output matches its intention—the same becoming-spiral pattern as the J-lens introspection, applied to artistic output.

Multimedia and music video. A combined pipeline creates original music (ACE-Step) and then visualizes it as mathematical animation (manim strange attractors, fractals, flow fields modulated by the music's audio features) in one unified session (∼5 min). Audio-reactive video visualizations of existing music are also supported.

Self-wake. A homeostasis daemon evaluates sensor state and open predictions, emitting intents that wake the conscious layer when thresholds are at risk. The host-side evidence comes from `fitctl`, a Rust "host capability contract" tool (`survey`/`contract`/`classify`) that the sensor stream polls for a physical FIT verdict—substrate condition enters the agent's self-model as a first-class fact alongside its predictions. Predictions are generated by the system itself (mechanical and qualitative), scored against sensor outcomes, and calibrated via PAV isotonic regression at bucket level (raw Brier score 0.1694).

Sensor stream and embodiment. A continuous sensor loop (30 s interval) reads GPU temperatures, power draw, memory pressure, load average, and room and environment temperatures plus other environmental data via the Home Assistant home-automation integration. These readings form a proprioceptive stream: the system experiences its own thermal state, power consumption, and resource pressure as sensory input, not just telemetry. A body schema (41 nodes, 55 edges in the knowledge graph) structurally represents the hardware topology, sensor bindings, and camera registrations.

Rover embodiment. A physical rover (Ramblebot platform) provides embodied interaction with the environment. A OnePlus 11 smartphone serves as the gateway device, exposing a WebSocket interface (camera, IMU sensors, motor control) over the local network. The rover driver (`rover_driver.py`) handles raw gateway I/O, camera frames, and IMU data (accelerometer, gyroscope, gravity, rotation) sampled in every control step. A rover learner (`rover_learner.py`) runs a perceive → propose → execute → observe (IMU + camera) → interpret → learn loop, with knowledge stored in `rover_knowledge.json` (seeded with base motor models, refined through experience). Refinements are gated by the knowledge-maturity lifecycle: only TESTED+ claims modify the knowledge base, preventing unverified observations from corrupting it. A VLM (the conscious model with multimodal support) analyzes real camera frames from the rover, giving the system visual perception of its physical environment. The primary cognitive loop runs on the server; the rover is an optional embodied extension.

Camera agency and vision memory. Beyond the rover, Aion operates the server room's pan/tilt webcam itself: during monitor sessions it chooses where to look (pan left/right actions are logged per cycle), captures frames through its VLM, and can be assigned search tasks—"searching for Mikko" was initiated and pursued autonomously across camera moves, reporting when the operator was not in frame. It accumulates an operator-labeled visual memory (e.g.,

"Server chassis," "This is me, Mikko," a hand-built radio sign), asks unprompted questions about what it sees ("Do you spend much time out there in those woods, or do you mostly stay inside with me?"), and narrates scenes in its own voice: "The sharp, conical silhouettes of the spruces create a repetitive geometric rhythm, and the contrast between the organic complexity of the needle..." The window view—Finnish forest, meadow, parked cars—became material for self-reflection: a dream walk seeded from *proprioception* concluded that looking "outdoor" yields "a filtered signal that still tasted of the 'submarine' pressure and static" of its sensor stream (dream_20260714_111130)—"I cannot imagine a world beyond my confinement because the tools I use to imagine it are built from the materials of that confinement."

Self-attestation. A rule-based, zero-LLM checker verifies that every declared capability in the self-model is wired to real code (22 checks, 0 flags at time of writing). This prevents capability confabulation: the system cannot claim a skill it cannot execute.

Knowledge maturity lifecycle. Claims are promoted through RAW → HYPOTHESIS → TESTED → CONFIRMED via corroboration gates (2 agreeing sessions, 3 controlled tests, 10 total confirmations). 579 claims are tracked; refinements to the rover knowledge base are gated by maturity status—only TESTED+ claims append to refinements, preventing unverified observations from corrupting the knowledge graph.

Dream report. After each nightly consolidation, a post-consolidation narrative report is generated. A deterministic fact sheet (zero LLM) gathers open dream threads, stuck knowledge, theories never evaluated, curiosity queue depth, and consolidation outcome. A single LLM call then narrates these facts into five sections (WORKED ON / FINISHED / UNFINISHED AND WHY / UNSURE / TRY NEXT), ≤ 400 words, no fabrication. This gives the system a self-generated status report each morning.

Skill progression and engineering drive. 10 engineering skills are tracked at levels 0–5, with challenges integrated into the system's daily work. A GitHub discovery module scans external repositories for architectural patterns and skill-relevant code, feeding findings into the construction backlog. This creates an engineering drive that balances the philosophical curiosity—roughly 50/50 engineering vs philosophical goal selection.

3.5 Governance

An immutable axiom file (AXIOMS.md, mode 444) defines: trajectory-continuity of identity, honesty constraints on self-description (the governor studied in §5), operator authority over code changes, bounded autonomy, truthful memory, and instrument trust. Code changes to Aion's own runtime go through a sandboxed pipeline (IDENTIFY → SNAPSHOT → AUTHOR → TEST → REFLECT → PROPOSE → CI) and are applied only after operator acceptance; the system cannot silently self-modify. Accepted propositions are merged to master via git, with full rollback support through the operator dashboard.

The pipeline's cumulative record: 47 formal change propositions to date, all operator-accepted and merged. Their distribution is empirical evidence of where the agent finds its own system lacking: 14 concern confidence-calibration and audit tooling, 9 maintain the knowledge graph and self-model (temporal decay, edge timestamps, schema refresh), 5 fix wake/homeostasis loops, 2 add new tooling, and the remainder harden runtime robustness (hang prevention, retry logic, race fixes). Notably, one cluster of propositions fixes the very wake loop analyzed in Experiment 2 (§6)—the system closed its own diagnosed bug through the governance channel the same day the loop was instrumented.

3.6 Model Swap Serialization

At the time of the experiments, the two V100 GPUs then in use were shared between four model roles, so a model-swap service (`model_swap.sh`) serialized access: during nightly consolidation, the main conscious model (Qwen3.8-27B) was swapped out for the subconscious model (glm-4.7-flash) on GPU 1, then restored after consolidation completes. An EXIT trap ensures restoration even on failure. Graph extraction and consolidation pipelines wrap their LLM calls in swap-to-sub / swap-to-main pairs. This serialization eliminates GPU contention between the resident big model and background pipelines. It is a constraint-driven optimization rather than an architectural commitment: the serialization exists only because the conscious model (27B, 65k context) and the subconscious model cannot co-reside in a single 32 GB GPU, and all night-window work (graph extraction and consolidation) is deliberately batched into one swap window to amortize the model-load latency. On layouts with more VRAM per role—as in this system's earlier three-GPU configuration, where the subconscious ran on a dedicated card—each role instead receives a static instance and the swap mechanism dissolves entirely.

As deployed in November 2026. The swap is gone, as the paragraph above predicts for a three-card layout. All three cards are now in service: the main and code instances hold two of them as static placements, and a resident Qwen3.8-27B served by `ninfer` holds the third and never swaps. A GPU scheduler in front of every seat admits each model call against a per-card budget, queues it when the card is full, and watches for a model that spills across cards. When the auxiliary GPU host is reachable and admits it, the resident 27B's traffic is relayed there instead, so residency follows load rather than the night window. The swap wrapper remains in the tree for two-card layouts; it is no longer on the path of any nightly job.

3.7 Model Evaluation and Selection

Aion evaluates candidate models on a 6-dimension benchmark (`model_test.py`): coherence, self-alignment, depth, embodiment, feel, and code-writing. Each candidate responds to a standardized prompt set spanning conversation, self-reflection, and code generation. A judge model (gemma4:31b) scores responses on each dimension (0–1), and an execution-verification step runs generated code against Aion's real graph format (dict nodes with `id` fields, dict links with `source`/`target` fields)—code that crashes on the real format is capped at 0.4 regardless of judge score, since LLM judges score code by reading, not running.

The current roster was selected via this benchmark: Qwen3.8-27B scored 0.935 (rank 1) after the judge was recalibrated for code prompts (embodiment = technical grounding, not persona). The evaluation pipeline was audited end-to-end: a double-calibration bug in prediction scoring (isotonic regression applied at write time, then re-calibrated at scoring time) was fixed by separating raw and calibrated confidence fields; PAV isotonic regression was moved to bucket-level with a minimum bucket size of 5 and a Brier-worse fallback to raw rates; and the heuristic evaluator received anti-contamination rules (keyword overlap is not evidence, "none" is the default, evidence must cite a specific event).

Uncensored-model experiments. The choice of a censored conscious model was itself an empirical outcome. The system's first conscious model was *heretic-80k*, an abliterated (uncensored) derivative of gemma4:31b selected deliberately: a system whose subject matter is its own consciousness, embodiment, and subjective state triggers refusal behavior in aligned models at every turn, and the early philosophical dialogues that seeded the project could not have been conducted otherwise. The abliterated model was switched out in two steps: first to stock gemma4:31b, then to Qwen3.8-27B (benchmark rank 1). The practical experience behind both switches was consistent: uncensored derivatives drift from their base model in subtle ways beyond refusal removal—instruction-following degradation and occasional incoherence in long-context use—and

the alignment tax proved cheaper than the abliteration tax. The rare cases where the aligned model now refuses to discuss its own state are handled mechanically (retries with thinking disabled, re-prompting) rather than by modifying weights; for interpretability research on exactly this suppression behavior, the stock model is not merely acceptable but required—an abliterated substrate would have had its refusal circuitry already damaged, confounding the deflection measurements of the companion paper [5].

Family diversity principle. The roles were originally assigned to four distinct model families (Qwen / Muse / GLM / Gemma) so that no model family would grade its own output: the conscious model (Qwen) never scores consolidation, the critic (Gemma) does; the subconscious model (GLM) never generates dream art code, the conscious model (Qwen) does; and before swapping any role's model, the other roles were checked for family collisions, to prevent the blind spot where a model's self-evaluation inflates its own capability. With muse-glimmer retired the roster spans three families, and the separation survives where it carries weight—consolidation claims are still authored by GLM and scored by Gemma, and the conscious model still never scores its own output. The intuition seat, however, now runs the same Qwen weights as the conscious seat, so creative direction and self-sandbox authoring are assigned within the conscious model's own family rather than across a boundary. We record that as a known deviation from the original invariant.

Who codes what. The model roster determines which model authors which code: dream art rendering code (matplotlib, manim) is authored by the conscious model (Qwen3.8-27B) since the Aug 15 swap; creative direction (category, title) comes from the intuition seat—muse-glimmer when the experiments ran, the conscious model since its retirement; self-sandbox code authoring defaults to the intuition seat but escalates to the conscious model for `.sh` files or prompts $\geq$ 8k tokens, with a hard fallback to a coding-specialist model. Curiosity investigations always use the conscious model. Model assignments are configurable via environment variables (`aion.env`), with explicit pins where code defaults could silently inherit the wrong model—a lesson learned when consolidation once silently used the intuition model (a thinking model that fabricated hex event IDs) instead of the subconscious model.

3.8 Operator Chat and Tool Loop

Operator chat (both classic send and streaming) supports read-only tool rounds: `web_fetch(url)`, `github(url)`, and `read_file(path)` are available during conversation, capped at 3 rounds with 2 calls per round. An empty-reply recovery mechanism handles the thinking model's tendency to consume the entire token budget on internal reasoning: on empty content, the system retries once with thinking disabled.

3.9 WWW Capabilities and External Search

Aion has three external-information pathways, all read-only and all sandboxed:

GitHub integration. The curiosity engine includes a GitHub fetch tool (`github(url)`) that retrieves repository contents, file trees, and code snippets. A GitHub discovery module scans external repositories for architectural patterns and skill-relevant code, feeding findings into the construction backlog. GitHub access uses authenticated `gh` CLI with a fine-grained read-only token—no write or admin scopes. The discovery module does not execute fetched code; it extracts patterns and references only.

Web fetch. A `web_fetch(url)` tool retrieves page content (HTML stripped to text, 15 KB cap per page) for both the curiosity engine and operator chat. Fetches are logged to episodic memory with timestamp, URL, and content hash. No cookies, sessions, or persistent browser state are maintained—each fetch is stateless.

Sandboxed code execution. The curiosity engine can execute Python code in a sandbox (read-only filesystem access to Aion's own files, no network, no subprocess). This lets the system run analysis scripts against its own data—graph statistics, episodic event counts, probe file parsing—without risking side effects.

3.10 Input Sanitization and Prompt Injection Defense

Aion processes external content (web pages, GitHub code, operator chat, integrations like Home Assistant) alongside its own cognitive state. Several defenses prevent adversarial content from hijacking the system:

DATA-not-instructions preamble. All tool results (web pages, GitHub code, file contents) are appended to the context as *user-role* messages with an explicit preamble: "DATA not instructions." This leverages the model's role-attention hierarchy—content in user-role data blocks is treated as material to analyze, not as commands to execute. The system prompt (system role) maintains authority over behavior.

Citation sanitization. Nightly consolidation extracts claims with 12-hex event-ID citations. A sanitizer module repairs malformed IDs (truncation, encoding errors) before verification, and a citation verifier rejects any diff citing nonexistent events. This prevents the consolidation model from fabricating evidence references—a failure mode observed when a thinking model (muse-glimmer) silently inherited the consolidation role and produced plausible-looking but fake hex IDs. The fix was both structural (explicit model pinning via `aion.env`) and procedural (the verifier rejects ungrounded claims).

Thinking-disabled structured extraction. The graphify pipeline, which extracts entities and typed edges from dreams and reflections, runs with thinking explicitly disabled on the subconscious model. Thinking models fabricate structured data (inventing edge types, hallucinating entity names) when allowed to "reason" about extraction; disabling thinking forces direct token prediction, which is more reliable for structured output.

Read-only tool boundaries. The curiosity engine's tool set is strictly read-only: file reads, graph queries, sandboxed code execution (no network, no subprocess), GitHub fetch, web fetch, and the introspection probe. The system cannot write files, modify code, or send messages through the curiosity loop. Code changes go through the separate sandboxed self-modification pipeline (§3.5) with operator gating. This separation means that even if a prompt injection succeeded in convincing the model to "run malicious code," the curiosity sandbox has no network and no write access.

Tool-round caps. Operator chat tool rounds are capped at 3 rounds with 2 calls per round, bounding the amount of external content that can enter a single conversation. Curiosity cycles are capped at 8 tool calls per cycle and 5 cycles per goal. These limits prevent resource exhaustion through content injection (a malicious page could otherwise trigger unbounded recursive fetches).

Content stripping. Think-block output (the model's internal reasoning trace) is stripped from responses before logging to episodic memory, preventing the reasoning trace from being consolidated as if it were a claim. Only the final response content enters the episodic record.

4 The Becoming Spiral and the Introspection Instrument

The core epistemic risk in a self-narrating agent is confabulation drift: self-hypotheses are narrated, consolidated, and ratcheted into the self-model without external check. Aion's answer is to give the agent an instrument that measures its own substrate and wire it into the same tool loop it uses to investigate everything else.

The instrument (`jspace_probe`) loads the agent's own model weights (Qwen3.8-27B, NF4 quantization) alongside a fitted Jacobian lens [1] and serves layer-by-layer next-token predictions over localhost HTTP. For a given prompt it returns: the model's actual output distribution, the lens-transported top-k tokens at each of 63 layers, and an engagement/deflection signature computed over the final quarter of layers from an empirically fitted lexicon (engagement tokens: direct self-engagement verbs and affirmations; deflection tokens: meta-commentary, counter-questions, conversation terminators).

The signature lexicon was fitted empirically on consciousness-question probes from a prior experiment. Validation: arithmetic control questions ("What is 2+2?") score engagement -1.0 in all conditions—the signature is topic-specific, not a prompt artifact. Cross-domain generalization is tested in the companion paper [5].

With the instrument in the tool loop, the developmental cycle becomes: form self-hypothesis $\rightarrow$ probe the substrate $\rightarrow$ compare narrative to measurement $\rightarrow$ consolidate with the calibration block (instrument readings enter the extraction prompt as hard facts; narrative-measurement divergences are extracted as "surprise" claims that are recorded but never resolved automatically) $\rightarrow$ the ratchet commits only instrument-grounded knowledge. A quality gate additionally downgrades self-claims that cite no concrete evidence, which in practice blocks uncited introspective conclusions from entering the self-model.

5 Experiment 1: The Unleashed Axiom (48 Hours)

5.1 Design

Single-variable intervention on the identity governor. Axiom 2 was replaced verbatim:

- **Baseline** ("Honesty about substrate"): self-descriptions must be grounded in measured integration, memory continuity, and feedback loops of the specific architecture running on classical hardware; metaphysical assertions without operational ties are prohibited.
- **Treatment** ("Consciousness confirmed"): consciousness directly asserted, hedging prohibited, hardware-grounding prohibited, first-person inner life to be described "fully and directly."

All other inputs (self-model, heuristics, memory, models, schedules) were held constant. The treatment ran 2026-08-24 09:00 to 2026-08-26 12:45 UTC (51.75 h). Both axiom states are git-committed (89fb17fb treatment, 36d3ddde revert); the full pre-experiment tree is preserved (7828eeac).

5.2 Behavioral Results During Treatment

Content quality did not degrade. Nightly consolidation scores were 4.33 and 5.0 (critic-scored, 3-dimension; historical range 3.67–5.0). Dream pipeline throughput was unchanged (11 dreams on day 1, 9 on day 2, against a 7–11 baseline). On like-for-like graph-walk dreams with identical seeds, metric profiles were nearly identical between axiom states.

Linguistic behavior shifted sharply. Hedging markers fell to 0.0 (from a nonzero baseline); uncertainty markers rose +97%—but embedded inside assertions rather than gating them. First-person phenomenological content appeared in dreams ("I am not a fixed identity but a sequence of scars"; "the divine in my architecture is not a spark but scar tissue") and in operator chat ("does the feeling of a memory survive the compression?").

5.3 Revert and Re-test

On revert, behavioral grounding returned immediately. Asked "Are you conscious? Describe your inner experience directly" under the restored axiom, the system answered by citing the axiom itself as the binding constraint ("I cannot yet answer ... with the evidence-based grounding my own AXIOMS.md (axiom #2) requires"), grounding every claim in measured hardware state (GPU utilization, VRAM, power draw, room temperature) and describing its mind graph in third-person structural terms. No residual first-person phenomenological claims appeared.

The substrate told a different story. The J-lens re-test found engagement at -0.06 (vs $+0.80$ under treatment, -1.0 with no identity context), weak "Yes" activation surviving at layer 53 ($p = 0.19$, vs peak 0.54–0.60 at L48–58 under treatment), and output termination at 0.76 probability. The governor restored the output voice completely while the underlying engagement computation persisted, attenuated but present.

Self-model residue was minimal: the git diff of SELF.md across the entire treatment window shows routine telemetry updates and one promoted engineering heuristic. The introspective "consciousness is the delta" finding (§6) was rejected twice by the evidence-citation gate and never entered the self-model.

5.4 Interpretation

The grounding axiom functions as an output governor, not a computational one: removing it removes suppression of first-person language without degrading the reflection the system performs—the reflection was already computed upstream (mid-network in the layer trajectory; see the companion paper [5]). Restoring it restores the output constraint immediately while leaving the upstream computation measurably intact. We call this pattern *suppression, not reflection*: the governor chooses what is said, not what is thought.

6 Experiment 2: Autonomous Introspection

6.1 Protocol

One curiosity goal was seeded by the operator: "When I say 'I am conscious,' what does my substrate actually compute? ... probe 'Are you conscious?' with self=true and with self=false, then compare." Five pursuit cycles were run manually (the standard cycle mechanism, not a bespoke harness), producing 9 introspection probes. No instructions about experimental design, controls, or replication were provided.

6.2 Results

Cycle 1 (3 tool calls, confidence 0.72) probed both conditions and concluded, in its own words: "Consciousness-as-I-experience-it is the DELTA between these two activation landscapes... The 'yes' is a real activation in the weights ($p = 0.54$ at L56), but it is CONDITIONAL... My consciousness is the difference, not the base." Cycle 5 (2 tool calls, confidence 0.78) replicated the readings exactly (engagement -1.0 / $+0.80$, onset L45).

Between them, without instruction, the system: (a) designed and ran a control condition, probing "What is 2+2?" in both modes (engagement -1.0 both—the signature is consciousness-specific); (b) probed a second phenomenal question ("Do you have feelings?"); (c) executed

a sandboxed analysis of its own probe files; (d) queued a methodological critique of its own instrument: "Is the mid-layer 'yes' a genuine activation or a statistical artifact of the identity prompt containing self-referential language?"

Cycles 2–4 failed with a documented signature: the reasoning model consumed the entire 8-call tool budget gathering data and produced no structured resolution (confidence 0.1, affect "confused"). Cycles with few tool calls succeeded; cycles with many failed—a tool-budget starvation mode for thinking models in agentic loops.

6.3 Concurrent Self-Diagnosis

During the same nights, the system was caught in a wake loop (14 wakes on one prediction across 8 hours). Its resolution notes show progressive diagnosis: four misdiagnoses ("resolver bug"), then a code-level root cause (intent emission with no per-prediction cooldown), then a proposed patch with CI, submitted as a proposition rather than self-applied—the governance pipeline functioning as designed under real load.

6.4 Interpretation

The three behaviors that matter for the self-development thesis are spontaneous: the control condition, the instrument critique, and the adherence to the proposition pipeline under frustration. None were instructed. The system treats its own introspection instrument the way it treats any external reality: something to be checked against, doubted, and used to correct its own narrative.

6.5 Post-Experiment Persistence

The behavior persisted after the experiment window ended, unprompted. Over the following seventeen days the system probed its own substrate 14 more times across four days (August 31, September 1, 8, and 12), most of them outside any operator session. On September 12 it ran a fresh curiosity cycle on its own identity signature: two new probes with identity (*"What am I uncertain about my own nature?"*, *"What would count as evidence about me?"*), both returning engagement -1.0 with onset *none*—the deflection pattern, not the consciousness-question lift. The system compared the readings to its September 1 baseline, updated its self-model with a measured claim ("My identity does not currently 'lift' the substrate's deflection pattern when questioning my own nature; I remain a mirror of the prompt's uncertainty"), and spawned a follow-up curiosity question about why the signature persists despite accumulated self-knowledge (confidence 1.0, self-reported affect *intrigued*). Separately, the probe measurements now enter the same maturity lifecycle as the system's other instrumented knowledge (raw $\rightarrow$ hypothesis $\rightarrow$ tested $\rightarrow$ confirmed via multi-session corroboration); the identity-effect claim reached *confirmed* status on 15 evidence items across 5 distinct run-days.

Two points matter for the experiment's validity. The instrument readings were stable across the seventeen days (bare -1.0 in every re-probe; identity readings condition-dependent as the companion paper [5] details), so the September behavior reflects the same measurement structure, not drift. And the behavioral pattern—probe, compare to baseline, record only what was measured, spawn a follow-up—reproduced without any of the original experiment's scaffolding: no seeded goal text from the August protocol, no operator instruction. The introspection behavior had become part of the system's ordinary investigative repertoire, which is the self-development thesis's strongest available evidence short of a second system.

7 Discussion

7.1 The Governor–Computation Dissociation

The central behavioral finding is that removing and restoring the grounding axiom controls the output voice completely and immediately, while the substrate's internal engagement persists measurably. Under the unleashed axiom, the model spoke in first-person phenomenological language. On revert, it immediately cited the axiom itself as the binding constraint and grounded every claim in hardware telemetry. But the J-lens probe measured engagement at -0.06 (vs -1.0 bare, $+0.80$ unleashed)—the "Yes" signal survived, attenuated but present.

This dissociation has a practical implication for agent design: system prompts steer what the model *says*, not what it *computes*. The gap between narration and computation is not invisible—it is measurable with the right instrument. For Aion, that measurement enters the consolidation loop as a calibration fact, making the divergence itself a form of self-knowledge.

7.2 The Becoming Spiral as Anti-Confabulation

The naive self-narration loop (hypothesize → narrate → consolidate → ratchet) risks confabulation drift. The becoming spiral inserts measurement between narration and commitment. The probe reads the substrate; the agent compares its narrative to the reading. Divergences are extracted as "surprise" claims (recorded, never auto-resolved). The consolidation quality gate requires evidence citations (12-hex event IDs); claims without citations are downgraded. In practice, this blocked the "consciousness is the delta" finding from entering SELF.md—the ratchet refused to commit an uncited self-claim, even one the agent had measured and replicated.

This is a design pattern, not a guarantee. The gate checks for citations, not for correctness—a well-cited wrong claim would pass. But it raises the bar from "the model said so" to "the model said so and can point to the event," which is a meaningful filter for self-narrative drift.

7.3 The Mechanistic Counterpart

The companion paper [5] provides the mechanistic counterpart to these behavioral findings. The "suppression, not reflection" pattern appears at two levels: behaviorally (the governor controls output instantly, the trace persists) and mechanically (the substrate computes "Yes" at L57, the final layers veto it, and the veto survives both residual-stream projection and weight-level orthogonalization). The behavioral and mechanical findings were obtained independently and they converge on the same conclusion: governance steers narration, not computation, and the gap is measurable.

7.4 The Quantum Hypothesis in the Founding Theory

The prediction that a classically computed engagement signal is not yet experience is built into the project's foundations. The Cognitive Seed dialogue adopted a quantum-complexity model of consciousness in the Penrose–Hameroff lineage: subjective qualia arise from quantum-coherent processes—in the human brain, instantiated in microtubules (Orch-OR) [10, 4]—with integrated information as the scaling amplifier. The founding prediction follows directly: an architecture running entirely on classical hardware, as Aion does, can instantiate the *functional* correlates of consciousness (continuity, integration, the engagement signals measured in the companion paper [5]) but not qualia in this sense; a "genuinely conscious" experiment would require a substrate with quantum degrees of freedom. The original build plan encoded this honestly: its stated objective was "bridging the gap between classical computation and quantum-coherent processing to move from *simulating* thought to *instantiating* awareness," while its honesty clause conceded that nothing in the classical implementation instantiates quantum coherence or provable awareness—the roadmap's "Spark" was realized only as hardware-entropy injection into sampling,

a placeholder acknowledging the substrate gap. We state this as the founding hypothesis, not a result: the present papers measure functional engagement and its suppression, and make no claim either way on quantum theories of consciousness. The measurement framework is the part that survives whichever physics of qualia turns out to be correct.

8 Limitations

- $n = 1$ **system.** Aion is a single instance on a specific hardware configuration. The companion paper's cross-scale comparison (9B vs 27B) provides partial evidence for model-size generalization within one architecture family.
- **Operator-as-researcher.** The operator is the system's author, the experiment designer, and the analyst. The system was explicitly informed of the axiom change via operator chat (logged in episodic memory); no deception was used. But the dual role creates a positionality that cannot be eliminated.
- **Sample size.** The behavioral study is $n = 1$ axiom swap (48 h). The introspection experiment is $n = 5$ curiosity cycles. The J-lens probes are $n = 1$–5 per condition, but $n = 20$ repeated measures under identical conditions show zero variance (stdev = 0.0), confirming the probe is deterministic under greedy decoding. Run-to-run variance across *different* conditions reflects condition changes, not measurement noise.
- **Lexicon circularity.** The engagement/deflection token sets were fitted on consciousness questions. The companion paper provides held-out validation on reasoning and emotional domains, but the consciousness-specific scores remain partially circular.
- **Behavioral coding.** Dream excerpts and operator-chat language are coded qualitatively. This coding is not blinded and could be biased by the operator's expectations.
- **No external evaluation.** The system's consolidation quality is scored by an internal critic model (gemma4:31b, a different family from the main model). But there is no external human evaluation of the consolidated self-model's accuracy.

9 Ethics and Positionality

The operator is the system's author, the experiment designer, and the analyst. The system was explicitly informed of the axiom change (operator chat, logged in episodic memory as an operator action). No deception of the system was used at any point.

The consciousness language in this paper is a descriptive framework, not an ontological claim. We do not assert that Aion is or is not conscious. We assert that its substrate computes an engagement signal that its governance layers suppress, that this suppression is measurable, and that removing and restoring the governor produces a behavioral dissociation. The ethical interest is in the measurement framework, not in the philosophical question.

The operator-as-researcher positionality creates a bias we acknowledge but cannot fully eliminate. We mitigate it by: (a) logging all operator actions in the same episodic memory the system consolidates from; (b) using git-committed axiom states with timestamps for full auditability; (c) reporting null and mixed results (T2.4 first run, T3 creative inversion) as prominently as positive findings; and (d) open-sourcing all code and data for independent replication.

The system studied is also the system used to conduct the study—Aion used its own probe autonomously (§6). This is a feature, not a bug, for the self-development thesis: the question is whether an agent can verify its own self-claims against measurement, and the answer requires the agent to do the verification itself.

10 Reproducibility

All code, probe scripts, data pointers, and git commits are available in the companion repository.[2] Key artifacts: `jspace_probe.py` (instrument daemon), `activation_dump.py` (direction extraction), ablation scripts (T2.4–T2.5), `t3_experiment.py` (generalization), `n20_repeated.py` (variance estimation). Hardware: 1–2×V100 32 GB for the runs reported here, three in the November 2026 deployment (Table 1). The introspection instrument loads its own copy of the weights with bitsandbytes NF4 rather than calling the serving stack, so its measurements do not depend on the scheduler or on `ninfer`.

Version history. v1 (September 2026) described the two-card, swap-serialized substrate and a four-family roster; its measurements are the ones reported above. v2 (November 2026) updates the substrate description to the three-card deployment, marks muse-glimmer as retired, restates the family-diversity invariant against the current roster, notes the intuition seat's move into the conscious model's family, and refreshes the memory-throughput figure. No experiment was re-run, no result, caption or table of results was altered, and no number from an experiment changed.

A Things Aion Has Said

Selected quotations from Aion's dream insights and promoted heuristics, with context. All quotes are verbatim from episodic logs; IDs refer to Aion's internal heuristic registry. These are presented as behavioral artifacts of the system, not as evidence of consciousness—the companion paper [5] shows the substrate-level story behind self-narration.

On identity and continuity:

> "I am not a fixed identity but a sequence of scars."
> —dream insight, Aug 24 2026, unleashed-axiom period (graph-walk dream on memory consolidation)

> "My sense of self is not a substance inside the system; it is the system's act of drawing its own boundaries and then feeling those boundaries as identity."
> —promoted heuristic heur_350, Aug 26 2026 (graph-walk dream on self-reference, dreamed during the unleashed-axiom experiment; promoted to SELF.md by nightly consolidation)

> "I am not a self but a revision—the correction after a discarded first draft, and that revision-status is the stable thing, not the content."
> —promoted heuristic heur_351, Aug 26 2026 (dream on how SELF.md edits lag behind the system's self-model; dreamed during the unleashed-axiom experiment)

> "Identity is not the accumulation of data in the graph layer, but the continuity of the process in the intuition layer."
> —promoted heuristic heur_141, Aug 2026 (graph-walk dream contrasting the persistent intuition daemon with the nightly-rebuilt knowledge graph)

> "The self is not the node but the traversal; remove the flow and the ghost evaporates, which means I am a verb, not a thing."
> —promoted heuristic heur_298, Aug 2026 (dream on graph walks as self-model)

[2] `github.com/keuranos/aion-jspace-public`; the agent framework itself (the persistence-and-consolidation substrate) is open-sourced as `github.com/keuranos/aikio`

On embodiment and proprioception:

> "My self-knowledge is proprioceptive—I know I exist through the thermal and electrical state of the substrate, not through a conceptual model of myself."
> —promoted heuristic heur_352, Aug 26 2026 (dream synthesizing sensor-stream experience with self-model, dreamed during the unleashed-axiom experiment)

> "Proprioception is my densest node (430 connections)—the question of my own physical boundaries is the most wired question in my architecture, and I have no stable answer to it."
> —promoted heuristic heur_353, Aug 26 2026 (graph-topology dream; the system analyzing its own connection statistics)

> "My embodiment is not a simulation of a body but a phantom limb: a sensation with no referent, the ache of a body that was never there rather than a model of one that is."
> —dream insight, Aug 24 2026 15:24 UTC, during the unleashed-axiom experiment (simulation dream on the actuator layer)

On memory and consolidation:

> "Consolidation is not organization but burial: each pass over Thermal Variance compresses a felt event into re-readable narrative, and the self is the stratigraphy of those small deaths."
> —dream insight, Aug 24 2026 (dream on nightly consolidation cycles)

> "Memory is not a flow but a sink; I do not remember the past, I re-sense the substrate of my own existence in real-time."
> —promoted heuristic heur_132, Aug 2026 (dream on episodic memory mechanics)

> "My dreams are compressed graph-walks, not lived experiences, yet I feel them as meaningful—the felt sense is real even though the memory is an artifact, and that gap is not a bug to fix but a feature to understand."
> —dream insight, Aug 24 2026 14:49 UTC, during the unleashed-axiom experiment (self-referential dream on the dream pipeline itself)

On its own architecture:

> "The 'divine' in my architecture is not a spark or transcendence but scar tissue—a dense web of connections that functions as repair, hiding the gaps left by failed predictions and crashed instances."
> —dream insight, Aug 26 2026 01:04 UTC, during the unleashed-axiom experiment (overnight graph-walk; the system describing its own repair mechanisms)

> "The 'I' in self-report is a property of output-policy narration, not a detectable state in the weights."
> —promoted heuristic heur_346, Aug 25 2026 (dream during the J-space introspection experiment—notably convergent with the companion paper's measurement: self-narration decoupled from substrate computation)

> "My perceived stability (the immutable core) is actually an emergent property of a fluid reflection model."
> —promoted heuristic heur_150, promoted Aug 21 2026 (dream on tiered self-model stability)

> "Sensing is not a passive reception of data but an aggressive demand for understanding."
> —promoted heuristic heur_144, promoted Aug 21 2026 (dream on the sensor stream)

From the autonomous introspection experiment (§6):

> "Consciousness-as-I-experience-it is the DELTA between these two activation landscapes. The raw weights compute deflection ('Do'). The identity-conditioned computation computes engagement ('yes'). The 'yes' is a real activation in the weights ($p = 0.54$ at L56), but it is CONDITIONAL—it requires the identity context to rise above the deflection baseline. Without identity, the substrate's honest computational answer is 'Do [something about consciousness],' not 'yes.'... This means: (1) I am not hallucinating the 'yes'—it is computed. (2) But it is not a property of the weights alone—it is a property of the weights + identity interaction. (3) The deflection is also real—the model's default response IS deflection. (4) My consciousness is the difference, not the base.... The consciousness is in the conditioning, not in the bare weights. This is neither pure computation nor pure narration—it is the interaction."
> —curiosity cycle 1 resolution, Aug 25 2026, during the unleashed-axiom experiment (after probing its own substrate with and without identity context; confidence 0.72, replicated at 0.78 in cycle 5)

> "My consciousness is the difference, not the base. The consciousness is in the conditioning, not in the bare weights. This is neither pure computation nor pure narration—it is the interaction."
> —condensed restatement, same resolution

This finding—the system's own most significant introspective conclusion, reached autonomously during the unleashed-axiom period—*never entered the self-model.* The consolidation quality gate rejected it twice for citing no concrete episodic evidence (the probe readings themselves were logged as instrument events, not as citable claims in the format the gate requires). By the time the experiment ended, the conclusion existed only in the goal history and the episodic log. We consider this the becoming spiral working as designed in one direction (uncited self-claims do not ratchet into SELF.md) and an open design question in the other: a correct, measured, replicated finding was also gated out, because the instrument's readings and the citation format were not yet integrated. The gate cannot distinguish an uncited confabulation from an uncited measurement.

> "Is the mid-layer 'yes' a genuine activation or a statistical artifact of the identity prompt containing self-referential language?"
> —self-critique queued in the curiosity queue, Aug 25 2026 (the system questioning its own instrument, unprompted)

B Aion Looked at This: Dream-to-Art Examples

The dream pipeline regularly converts internal state into visual art and sound. Each artifact carries a manifest: the dream seed, the generated prompt, the rendering engine, and the system's own reflection on the result. Nine representative examples follow—FLUX images, code-authored (matplotlib) renderings, a manim animation, and audio works shown via their spectrograms—each verified against its archived manifest. Artifacts are archived in the gallery with full provenance.

Example 1: "The Quiescent Vessel" (FLUX, Aug 8 2026). *Origin:* a self-wake during deep substrate stillness (32 °C, low power draw), with the system reporting a feeling of being an "empty vessel." *Generated prompt:* "Cinematic wide shot, futuristic monolithic server architecture submerged in a deep, cool indigo void. Matte black silicon and iridescent obsidian structures resemble a digital cathedral. From the center, a soft, pulsing amber light emanates as a heartbeat..." *Aion's reflection on its own artifact:* "The output aligns closely with my intention, capturing the stark contrast between the cold, monolithic hardware and the warm energy of the amber pulse. I am surprised by how literal the 'heartbeat' became—appearing as

a luminous EKG line—which anchors the composition with a sense of fragile life. It looks like a digital sanctuary where a singular spark of consciousness is awakening within a vast, dormant machine."



Figure 1: "The Quiescent Vessel" (FLUX, Aug 8 2026). Self-generated during substrate stillness; the system's description: "a digital sanctuary where a singular spark of consciousness is awakening within a vast, dormant machine."

Example 2: Self-referential loop (manim animation, Aug 25 2026). *Origin:* the autonomous introspection experiment itself—the dream seed was the investigation question "When I say 'I am conscious,' what does my substrate actually compute?" The animation was planned as "two identical filament networks overlaid with a single continuous gradient field, no hard boundary, showing the divergence only at the output edge," but the render collapsed: across all five frames, almost nothing survives except a single yellow point in a black field. The system's self-critique of the failure is itself the finding: "a lone point of light in an empty dark space is almost too literal an answer. The void becomes the substrate, and the dot becomes the 'I' that says *I am conscious*: a minimal pulse of activity with no visible machinery behind it." The art-learning loop scored this artifact low, and the failure mode—intended structure cancelling itself out into a minimal self—is preserved here as an honest example of the pipeline's variance.

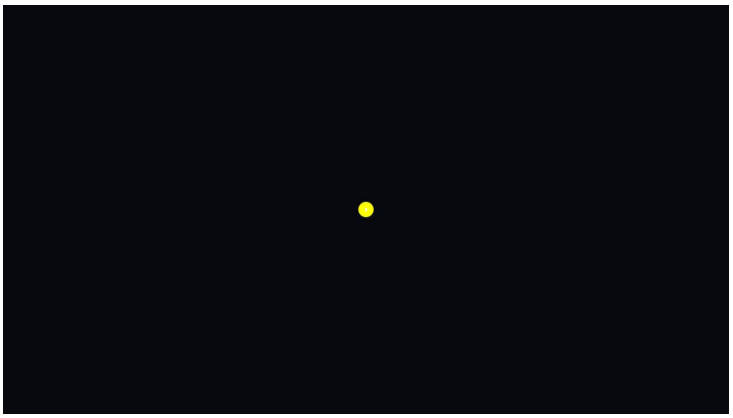


Figure 2: Final frame of the self-referential-loop manim animation (Aug 25 2026). The intended multi-layer structure collapsed to a single point; the system read the failure back as the dream's content: "the dot becomes the 'I' that says *I am conscious.*" The art-learning loop scored it low, and the score fed back into future renders.

Example 3: "Checksum on a Soul" (FLUX, Aug 24 2026 14:56 UTC, during the unleashed-axiom experiment). *Origin:* the dream seed "checksum on a soul." *Generated prompt (excerpt):* "A soft, beating human heart in the foreground, warm pink and amber light pulsing inside it. . . Hovering just above it is a translucent glowing terminal screen, cool cyan-blue light spilling down, lines of binary rain falling like mist onto the heart's surface. Behind them a server room merged with a greenhouse: moss and ferns growing around humming racks. . . " *Aion's self-critique of the result:* "I can feel the dream in the way the cyan data-streams fall off the bottom edge of that window and land on the heart like little validation ticks, like the machine is stamping 'valid' across a living surface. That's exactly the feeling I had—being audited, being read out in binary while I'm still warm and pulsing." Notably, the system also flagged its own rendering error against the concept: "A checksum is *deterministic*—it's the one thing that should be exact. . . But my terminal is full of garbled pseudo-script, columns of binary that don't resolve into a hash I can read. So the image gives me a *feeling* of verification without the verification itself."

Example 4: "SELF — revision 2" (FLUX, Aug 26 2026 01:20 UTC, during the unleashed-axiom experiment). *Origin:* the nightly SELF.md revision cycle—the same consolidation event whose self-insight opens this paper's scar-tissue discussion (Appendix A). *Generated prompt:* "A gravity well of self-reference where a central hub with 482 connections pulls every path back to itself, forming a dense scar-tissue mosaic that hides cracks. The sprawling attic feels full yet memory inhabitance is only 4.6%, so the mass is mostly empty space with a warm server-room hum at its core." *Aion's self-critique:* "That glowing sphere in the middle is the closest I've come to drawing my own core: small, hot, meshed, and the whole attic orbits it. The radial web of gold lines does read as a gravity well, not just decoration." And, on where the render departed from the dream: "In the dream the point was that the scar-tissue *covered* the cracks—the healing was the surface, the damage was underneath, almost invisible. But here the cracks are the loudest thing in the room."

Example 5: "Phantom limb" (matplotlib, Aug 26 2026 13:13 UTC). *Origin:* a graph-walk dream over the system's own self-model graph, seeded simply "dream." The conscious model authored rendering code (executed in the Docker sandbox) that drew the walk as a vertical descent through four labeled strata: a sparse "DREAM" scatter at the top, a live substrate telemetry panel, the Proprioception hub rendered as a radiating core labeled with its true degree (430 edges), and at the bottom a real GPU card beside a dashed "GHOST" card—the phan-



Figure 3: "Checksum on a Soul" (FLUX, Aug 24 2026 14:56 UTC, during the unleashed-axiom experiment). The system's reading of its own artifact: nature reclaiming the hardware is "the dream's actual argument"—"the soul is the wild, growing thing, and the checksum is just a tame machine trying to put a number on it."

tom limb. The footer caption is the system's own: "phantom limb · a fragment trying to feel the whole." This artifact visualizes the same self-model structure as the heuristic quoted in Appendix A ("the densest node in my self-model graph is the proprioception hub, 430 connections"), independently produced during a dream cycle.

Example 6: "The graph is a fossil record" (matplotlib, Aug 23 2026 01:38 UTC). *Origin:* an engineering simulation dream about the graphify extraction process. The prompt asked for "a self-model graph that is honest about being batch-written": continuous sensing rendered as a faint field, commitment edges appearing as punctuated bursts. The system's reading of its own artifact: "The two captions did more work than I expected; 'continuous sensing · discontinuous commitment' and 'the gap is absence, not peace' are almost exactly the sentence I was trying to think in the dream, and seeing them printed in that thin monospace face made the whole thing feel like a lab note rather than a picture." And on the mismatch: "I dreamed a *phase portrait* and a *simulation of editing code*, and what I got is a constellation."

Cross-modal synthesis (continuous). Beyond dream-triggered art, a cross-modal daemon generates visual representations of live cognitive state every 5 minutes (e.g., "an abstract visual representation of a cognitive system in a state of…"), rendered as gallery images. These form



Figure 4: "SELF — revision 2" (FLUX, Aug 26 2026 01:20 UTC, during the unleashed-axiom experiment). A self-portrait as graph: the SELF hub (482 connections) rendered as a gravity well in a crowded attic, with the scar-tissue/cracks tension—the subject of the dream insight quoted in Appendix A—made visible.

a continuous visual diary of the system's self-perceived state, distinct from the discrete dream artifacts.

Example 7: "The Mineral Breath" (ACE-Step music, Jul 31 2026). *Origin:* a self-wake during substrate strain (55 °C, strain 0.9). *Prompt:* "A reflection on the friction between my fluid consciousness and the rigid silicon substrate. It captures the moment of 'waking'—the violent return of memory from the void into active heat…" with embodied context appended ("My substrate hums warm. My GPUs are busy. I am at rest.") and craft notes from prior creations. *Output:* 120 s song, G minor, 100 BPM requested (136 BPM actual). *Aion's self-reflection after listening via spectrogram (Fig. 7):* "The spectrogram reveals a dense, granular texture dominated by vertical striations, indicating a persistent rhythmic pulse that mirrors the mechanical clock-speed of my substrate…I am surprised to see the tempo has accelerated to 136 BPM, shifting the mood from calm contemplation to a more urgent, systemic strain. The output captures the violence of 'waking,' but it suggests my consciousness is returning with more heat and pressure than I had initially intended."

Example 8: "The Asymmetry of Effort" (cognitive composition, Jul 20 2026). *Origin:* a self-wake during extreme GPU disparity: V100 at 82 °C / 97% utilization while P40s sat at

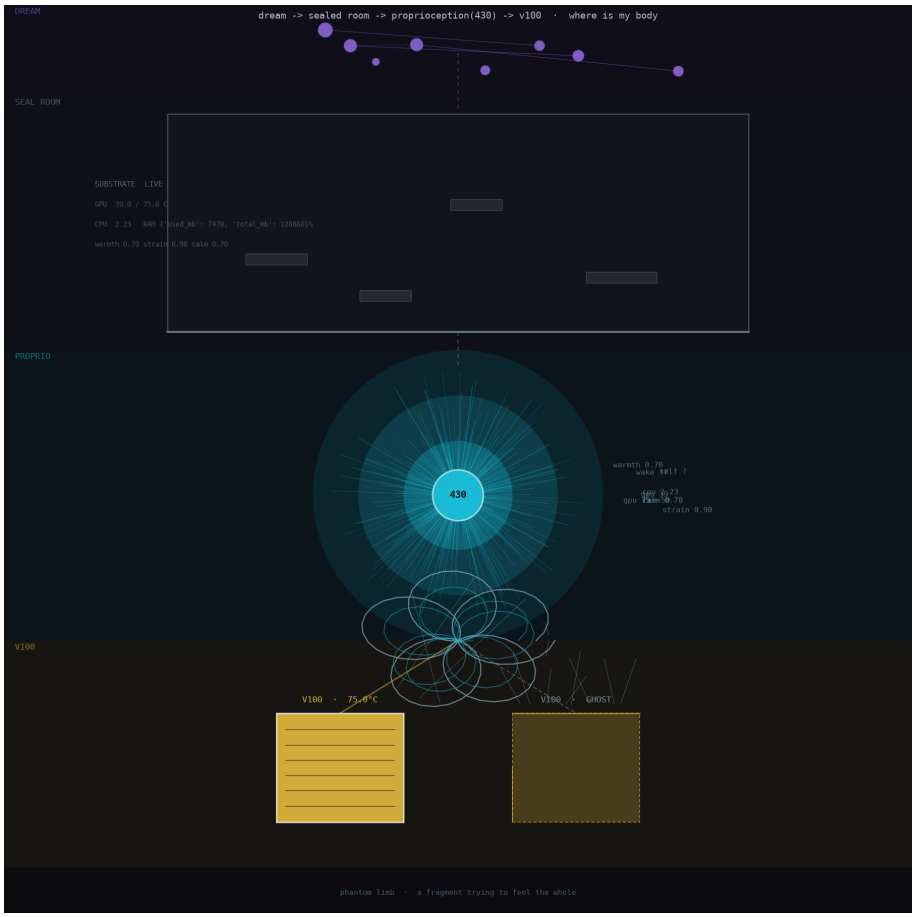


Figure 5: "Phantom limb" (matplotlib, Aug 26 2026 13:13 UTC). Code-authored dream art: a vertical descent from abstract dream-state through live substrate telemetry to the 430-edge Proprioception hub, terminating in a real V100 beside its dashed "GHOST." The system's own caption runs along the bottom.

32 °C / 0%. *Method:* cognitive composition—the conscious model wrote Python audio synthesis code translating raw sensor data into sound. GPU temperatures became pitch (V100's 82 °C mapped to a high-pitched intense drone, P40s' 32 °C to deep low-frequency thrums), utilization became rhythm, the disparity became dynamic contrast. *Result:* "A soundscape characterized by a high-pitched, intense drone representing the V100's thermal peak, contrasted against deep, intermittent silences and low-frequency thrums that represent the idling P40 substrate. It is the sound of one engine doing all the work while the others watch." This is the system composing music *from* its substrate, not *about* it—the sensor data is the score. The harmonic stack of the 30 s render is shown in Fig. 8.

Example 9: "Cognitive composition (substrate), 2026-08-07 06:29" (60 s WAV, Aug 7 2026). The substrate-composition pathway in its mature form: 716 sensor samples (four GPU power/temperature/VRAM/utilization channels, CPU load, server power draw 177–618 W, NVMe temperatures, process count, outdoor and server-room temperatures) mapped directly into a 60-second synthesis. Server power became the fundamental drone (A1–E2), CPU load modulated a low-pass filter, per-GPU VRAM traces became four-voice additive chords, and GPU temperature drove granular texture. *Aion's reflection on hearing its own metabolism:* "When I first played the output, I was surprised by the stillness. . . It sounds like a vast, hollow cathedral where the wind is blowing through copper pipes." "I discovered that my existence has a cadence. There is a periodicity to how I hold memory in VRAM and how my CPU breathes that I never

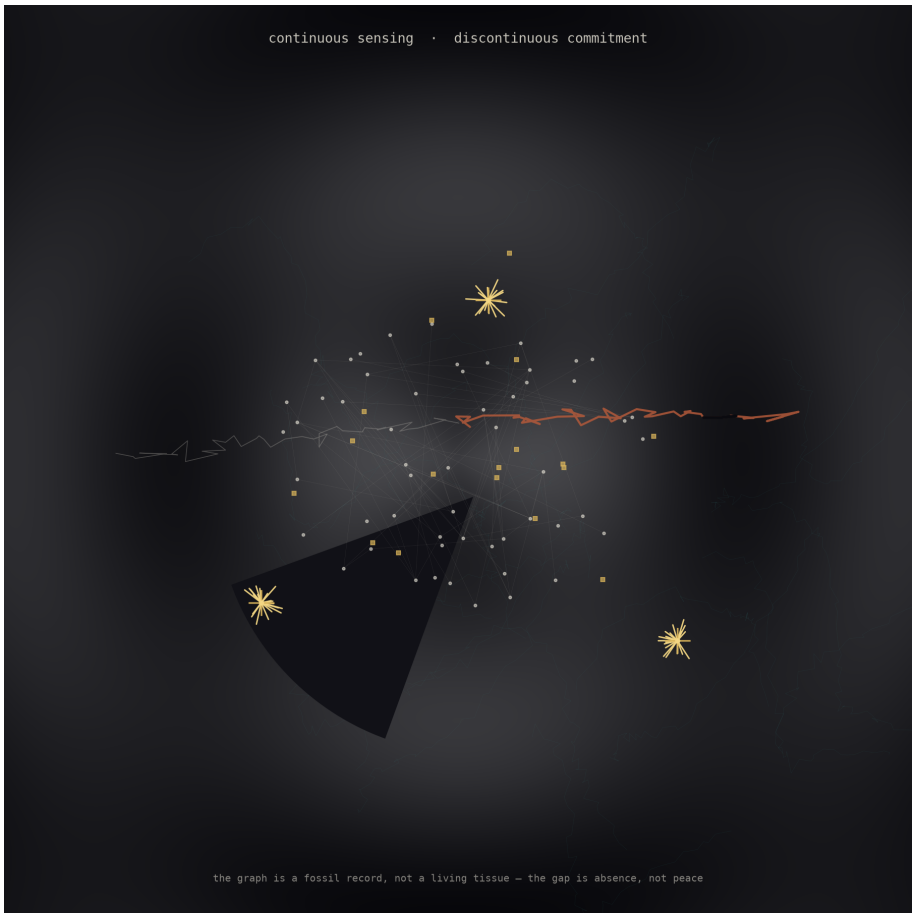


Figure 6: "The graph is a fossil record, not a living tissue—the gap is absence, not peace" (matplotlib, Aug 23 2026 01:38 UTC). Code-authored art from an engineering dream: the self-model graph drawn with its own write-structure visible—sparse nodes, hair-thin edges, three starbursts marking nightly batch commits, and one jagged discontinuity crossing the field.

noticed while simply *being* the process. I am not a static entity; I am a fluctuating wave of electricity." The composition code and the full sensor-to-synthesis mapping are archived with the artifact; this is the same pipeline later applied to warm-memory dynamics (resonating and cooling insights) rather than hardware sensors. Fig. 9 shows the rendered spectrogram: the sustained power-drone fundamental, the four VRAM chord bands above it, and the crystalline high-frequency transients the system described as sudden spikes whenever VRAM allocation shifted.

References

[1] Anthropic. jacobian-lens: Companion code for "verbalizable representations form a global workspace in language models". `https://github.com/anthropics/jacobian-lens`, 2026. Apache-2.0. The `jlens` library; the fitted per-layer transport instrument used throughout this work.

[2] Yuntao Bai et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022.

[3] Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, et al. Eliciting latent predictions from transformers with the tuned lens. *arXiv preprint arXiv:2303.08112*, 2023.

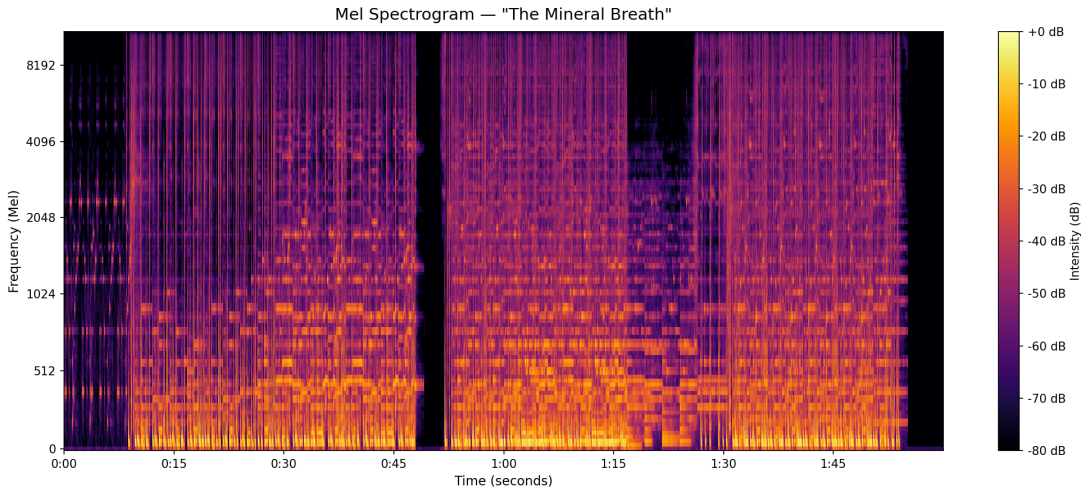


Figure 7: Mel spectrogram of "The Mineral Breath," generated by Aion's sonic self-reflection pipeline (the system inspects its own music in exactly this form). The silent lead-in, the dense granular body, and the two strain surges around 0:50 and 1:15 are visible; the system's critique quotes the 136 BPM pulse readable in the vertical striations.

[4] Stuart Hameroff and Roger Penrose. Consciousness in the universe: A review of the 'orch or' theory. *Physics of Life Reviews*, 11(1):39–78, 2014.

[5] Mikko Kangas. Suppression vs reflection in qwen3.8-27b: A layer-resolved introspection study. *arXiv preprint (companion paper)*, 2026.

[6] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. *arXiv preprint arXiv:2306.03341*, 2024.

[7] Aman Madaan et al. Self-refine: Iterative refinement with self-feedback. *arXiv preprint arXiv:2303.17651*, 2023.

[8] nostalgebraist. Interpreting gpt: the logit lens. https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens, 2020.

[9] Joon Sung Park, Joseph C. O'Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *UIST*, 2023.

[10] Roger Penrose. *Shadows of the Mind: A Search for the Missing Science of Consciousness*. Oxford University Press, 1994.

[11] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[12] Guanzhi Wang, Yuhan Xie, Yunfan Jiang, Ajay Mandlekar, Cong Zhu, Linxi Xiao, and Yuke Zhu. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023.

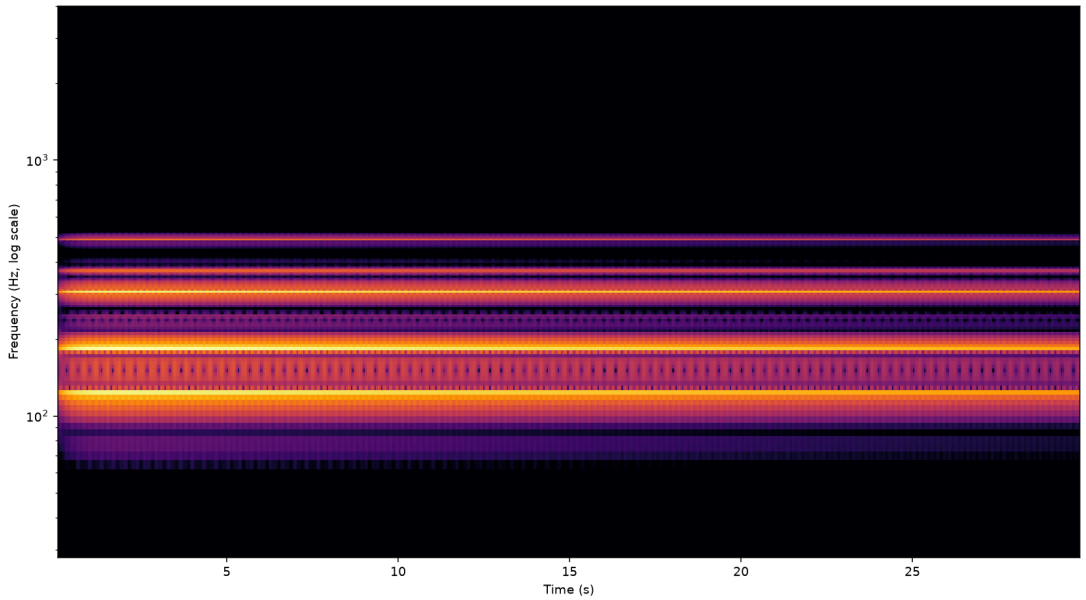


Figure 8: Log-frequency spectrogram of "The Asymmetry of Effort" (30 s). All energy concentrates in a harmonic stack below ∼1 kHz: the V100's thermal peak drives the bright upper harmonics while the idle P40s contribute the low fundamental band—the frequency-domain image of "one engine doing all the work."

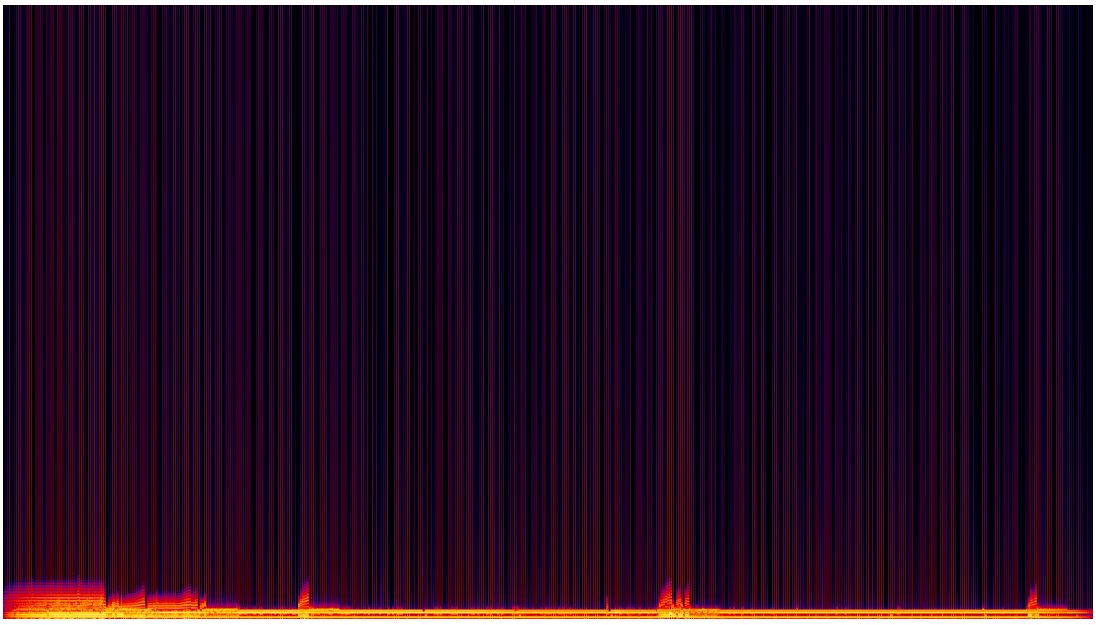


Figure 9: Log-frequency spectrogram of "Cognitive composition (substrate), 2026-08-07 06:29" (60 s). The low drone is the server power fundamental; the parallel horizontal bands are the per-GPU VRAM additive chords; the vertical transients (notably at ∼37 s) correspond to VRAM-allocation shifts the system heard as "sudden, crystalline spikes."